

# Chapter 2

## Intelligent Agents

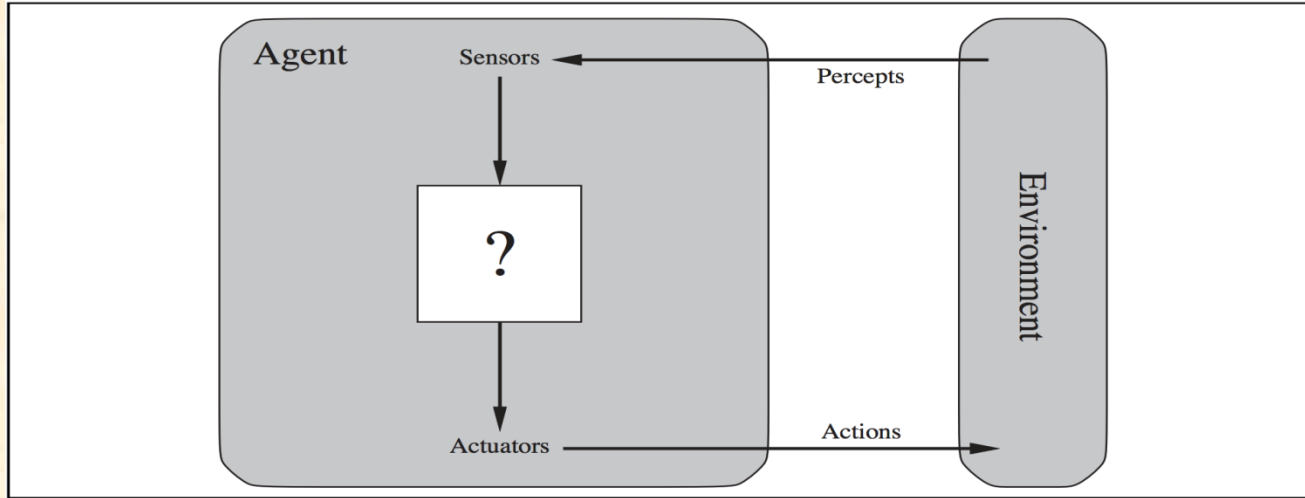


# Outline

- Agents and environments
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

# Agents

An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**



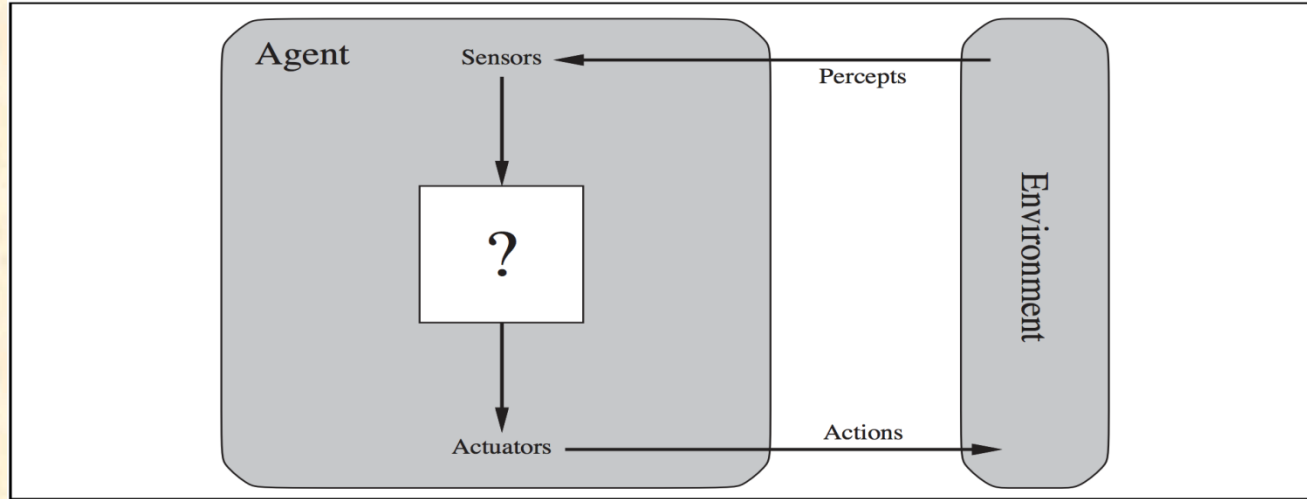
**Human agent:** **sensors** - eyes, ears, and other sensory organs; **actuators** - hands, legs, mouth etc

**Robotic agent:** **sensors** - cameras and infrared range finders ; **actuators** - various motors for actuators

**Software agent:** **sensors** - Key strokes, file contents, network packets  
**actuators** - Displaying on screen, writing files, sending network packets

# Agents

An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**



**Percepts:** agents perceptual inputs at any given instance

**Percept sequence:** complete history of agents perceived inputs


**Agent functions:** describes the agent's behavior

The **agent function** maps from percept histories to actions:


$$[f: \mathcal{P}^* \rightarrow \mathcal{A}]$$

The **agent program** runs on the physical **architecture** to produce  $f$

agent = architecture + program



# Vacuum-cleaner world



- Percepts: location and contents, e.g., [A, Dirty]
- Actions: *Left, Right, Suck, NoOp*

# A vacuum-cleaner agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

```
function REFLEX-VACUUM-AGENT([location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

What is the **right** function?

Can it be implemented in a small agent program?

# Good behavior: Concept of Rationality

## Rational agent

- An agent should strive to "do the right thing", based on what it can perceive and the actions it can perform.
- All entry in the table for the agent function are filled correctly.
- The right action is the one that will cause the agent to be most successful

## How to measure success???

- Performance measure: A criterion for success of an agent's behavior
- E.g., performance measure of vacuum cleaner
  - amount of dirt cleaned up,
  - amount of time taken,
  - amount of electricity consumed,
  - amount of noise generated, etc.

## Selection of performance measure:

- Average performance Vs energetic performance with long breaks
- Reckless life or safe but simple existence
- Economy with moderate poverty for all or some are super rich + some are very poor





# Good behavior: Concept of Rationality

## What is rational??

- The **Performance measure** that defines the criterion of success
- The agent's prior knowledge of **environment**
- The **actions** that the agent performs
- The agent's **percept sequence** to data

***Def: For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.***

How to check if Vacuum cleaner is a rational agent?

- **Performance measure:** award one point for each clean square
- **Environment:** Geography of environment is known apriori – two squares A & B
- **Actions:** Left, Right, Such and NoOp
- **Sensors:** Agent perceives its location and checks if contains dirt or not

Under these circumstances Vacuum cleaner agent is rational.

# Rational agents

- Agents can perform actions in order to modify future percepts so as to obtain useful information - **information gathering, exploration**
- An agent should **learn** from what it perceives
- An agent is **autonomous** if its behavior is determined by its own experience (with ability to learn and adapt)

# Nature of Environment

- **Building rational agents**
  - Task environments : “problems”
  - Rational agent : “solutions”
- **Specifying task environment**
  - PEAS: (Performance measure, Environment, Actuators, Sensors)
  - e.g., the task of designing an automated taxi driver:
    - **Performance measure**: Safe, fast, legal, comfortable trip, maximize profits
    - **Environment**: Roads, other traffic, pedestrians, customers
    - **Actuators**: Steering wheel, accelerator, brake, signal, horn
    - **Sensors**: Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

# PEAS

- Medical Diagnosis
- Part-picking robot

# PEAS

- **Agent: Medical diagnosis system**
  - **Performance measure:** Healthy patient, minimize costs, lawsuits
  - **Environment:** Patient, hospital, staff
  - **Actuators:** Screen display (questions, tests, diagnoses, treatments, referrals)
  - **Sensors:** Keyboard (entry of symptoms, findings, patient's answers)
- 
- **Agent: Part-picking robot**
  - **Performance measure:** Percentage of parts in correct bins
  - **Environment:** Conveyor belt with parts, bins
  - **Actuators:** Jointed arm and hand
  - **Sensors:** Camera, joint angle sensors

# Environment types

- **Fully observable vs. partially observable:**
  - Fully observable - An agent's sensors give it access to the complete state of the environment at each point in time.
  - Sensors detect all aspects relevant to choice of action (performance measure)
  - Agent need not maintain any internal state to keep track of the world
  - Partially observable - due to noise or inaccurate sensors or parts of state are missing
  - E.g. taxi cannot what other drivers are thinking
- **Deterministic vs. stochastic:**
  - Deterministic - The next state of the environment is completely determined by the current state and the action executed by the agent.
  - E.g. taxi driver is stochastic, vacuum cleaner example is deterministic
  - (If the environment is deterministic except for the actions of other agents, then the environment is **strategic**)
- **Episodic vs. sequential:**
  - The agent's experience is divided into atomic "episodes"
  - each episode consists of the agent perceiving and then performing a single action, and the choice of action in each episode depends only on the episode itself.
  - In sequential environment, the current decision could affect all future decision
  - E.g chess playing , taxi driver – sequential
  - Fault detecting robot in assembly line - episodic

# Environment types

- **Static vs. dynamic:**
  - The environment is changed while an agent is deliberating – dynamic else static.
  - Static – easy to deal, no importance to time
  - (The environment is **semi dynamic** if the environment itself does not change with the passage of time but the agent's performance score does)
  - E.g. taxi – dynamic, crossword puzzle – static, chess with clock - semidynamic
- **Discrete vs. continuous:**
  - Applied to state of environment, way time is handled, percepts/actions of agents
  - Discrete - A limited number of distinct, clearly defined percepts and actions.
  - E.g taxi – continuous ,
  - chess, crossword - discrete
- **Single agent vs. multi agent:**
  - Single - An agent operating by itself in an environment
  - E.g. crossword - single
  - Competitive multiagent – chess with opponent
  - Cooperative multiagent – taxi driving

# Environment types

	Taxi driving	Chess playing with clock	Chess playing without clock
Observable	Partially	Fully	Fully
Deterministic	Stochastic	Strategic	Strategic
Episodic	Sequential	Sequential	Sequential
Static	Dynamic	Semi-dynamic	Dynamic
Discrete	Continuous	Discrete	Discrete
Agents	Multi	Single	Single

The real world is partially observable, stochastic, sequential, dynamic, continuous, multi-agent



# Agent functions and programs

- An agent is completely specified by the agent function mapping percept sequences to actions
- One agent function (or a small equivalence class) is rational
- Aim: find a way to implement the rational agent function concisely

# Table-lookup agent

action  $\leftarrow$  LOOKUP( percepts, table)

- Drawbacks:
  - Huge table
  - Take a long time to build the table
  - No autonomy
  - Even with learning, need a long time to learn the table entries

# Agent program for a vacuum-cleaner agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

```
function REFLEX-VACUUM-AGENT( [location, status] ) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

What is the **right** function?

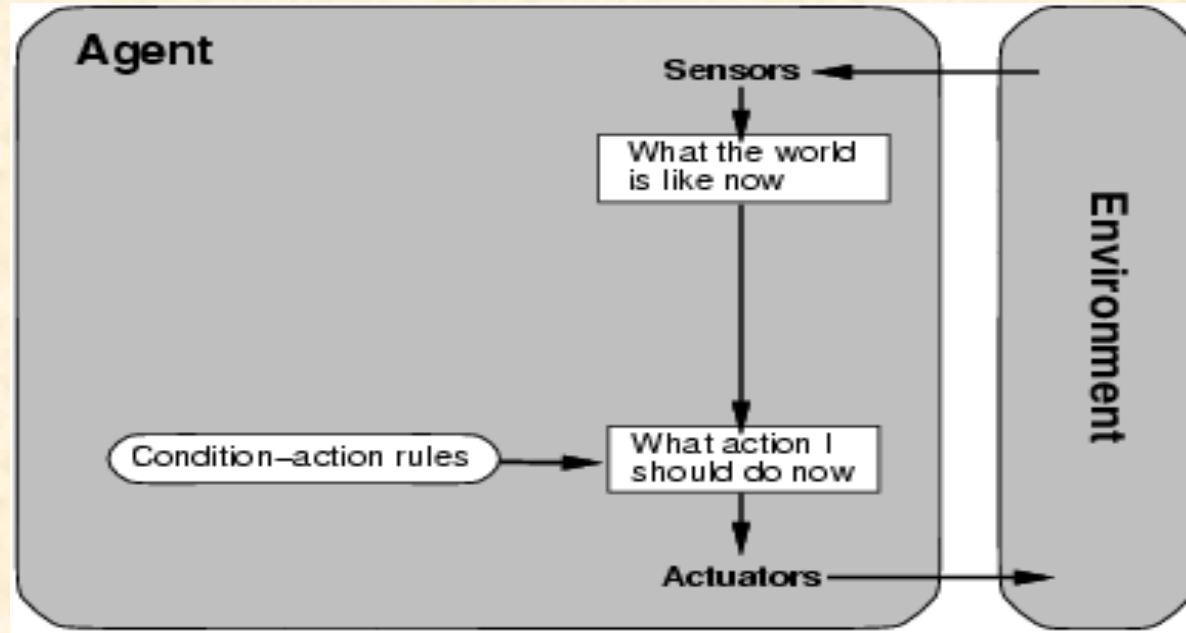
Can it be implemented in a small agent program?

# Agent types

Four basic types in order of increasing generality:

- Simple reflex agents
- Model-based reflex agents
- Goal-based agents
- Utility-based agents
- Learning agents

# Simple reflex agents



- Select action on current percept
- Ignore percept history
- Actions based on condition-action rule
- Problems?? – works only in fully observable environment

# Simple reflex agents

**function** SIMPLE-REFLEX-AGENT(*percept*) **returns** an action

**static:** *rules*, a set of condition–action rules

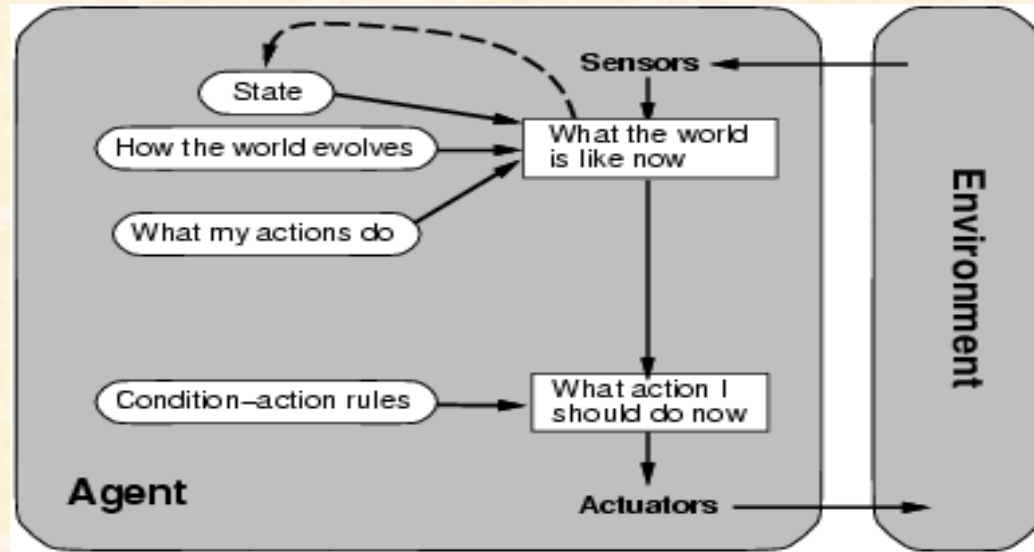
*state*  $\leftarrow$  INTERPRET-INPUT(*percept*)

*rule*  $\leftarrow$  RULE-MATCH(*state*, *rules*)

*action*  $\leftarrow$  RULE-ACTION[*rule*]

**return** *action*

# Model-based reflex agents



- Keep track of current state of the world agent can't see
- Maintain an internal state that depends on percept history n reflect unobserved aspects
- updating internal model requires knowledge about how the world evolves independent of the agent and how agent actions affect the world
- Model of the world → knowledge about how the world , hence the name



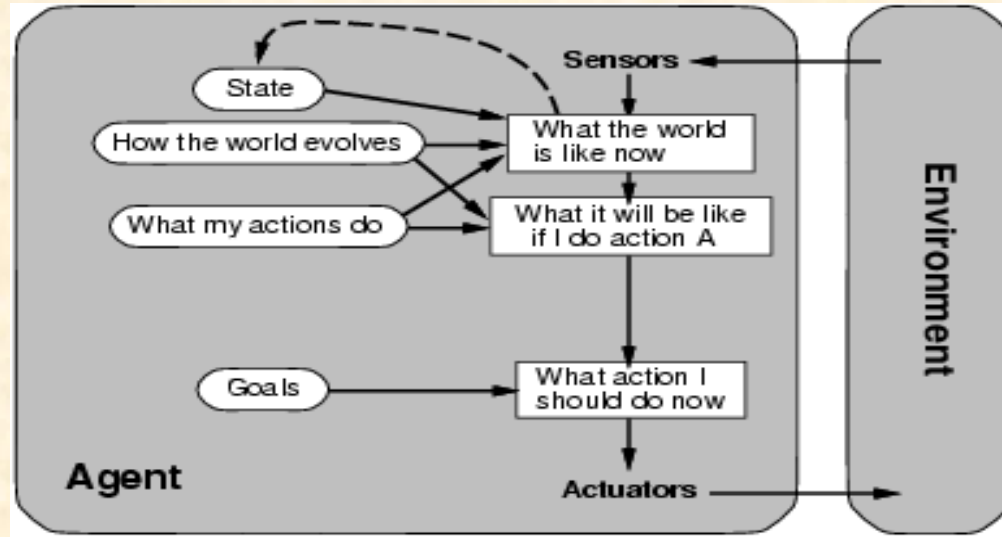
# Model-based reflex agents

```
function REFLEX-AGENT-WITH-STATE(percept) returns an action
  static: state, a description of the current world state
           rules, a set of condition–action rules
           action, the most recent action, initially none

  state ← UPDATE-STATE(state, action, percept)
  rule ← RULE-MATCH(state, rules)
  action ← RULE-ACTION[rule]
  return action
```

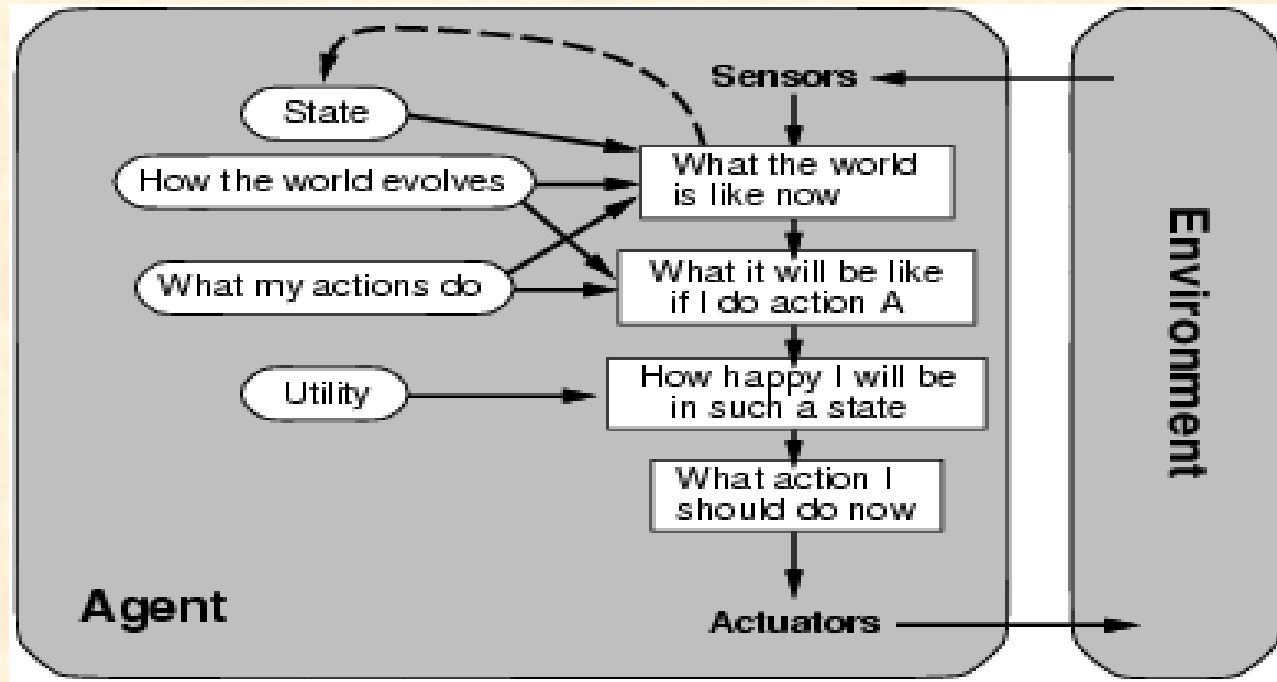


# Goal-based agents



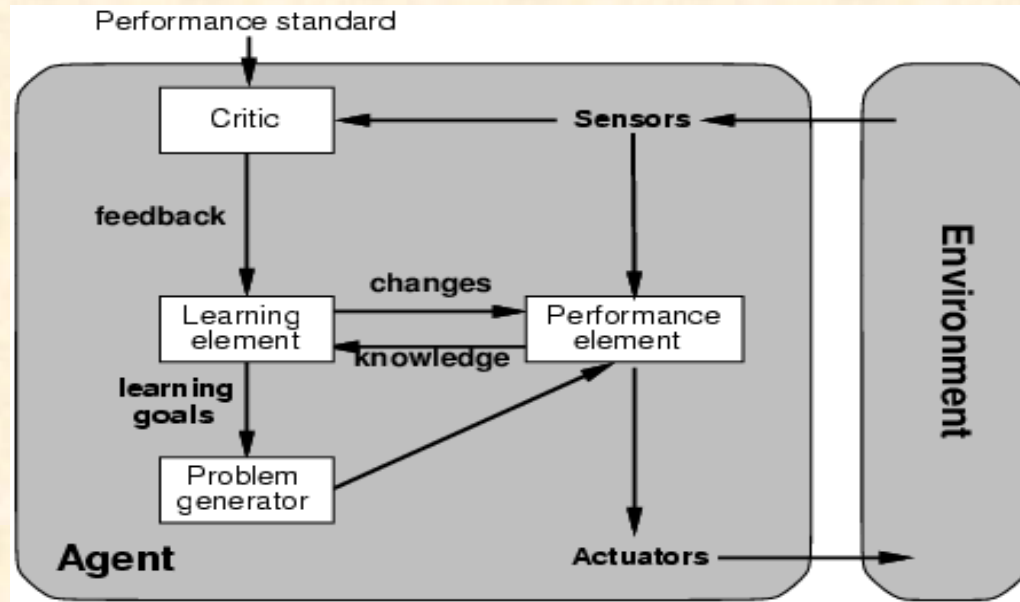
- Keep track of current state of the world and set of its goals
- Search and planning are required to achieve goals
- Decision making involves consideration of future – what will happen and will I be happy
- It is less efficient but more flexible s knowledge is explicit and can be modified
- Agents behaviour can be changed to a different location

# Utility-based agents



- Goals alone are insufficient to generate high quality behaviour
- Goals provide crude distinction between happy and unhappy state – need how happy
- Instead higher Utility of agent is preferred
- utility function maps a state to real number – degree of happiness
- Allows rational decision where goal is inadequate : in conflicting goals and in selecting goals that can be achieved

# Learning agents



- Learning allows agent to operate in initially unknown environment
- Learning element : making improvements
- Performance element : selecting external actions
- Critic : provides feedback on how the agent is doing and determines how the performance element should be modified
- Problem generator: suggesting actions

# Conclusion

- Agents can have variety of components
- Each component can be presented in a variety of ways
- A variety of learning methods but one unifying theme
- Learning in intelligent agent is a process of modification of each component to bring them into a closer agreement with available feedback information thereby improving overall agent performance